

Scalability in AI Research

By Manika Singh, Sabha Amrin, Gaurav Singh Parihar & Seema Verma

Abstract


Scalability plays a vital role in modern AI research, allowing models to make full use of powerful computing resources for faster training and better performance. In this report, we share our experience setting up, running, and evaluating distributed training experiments using GPUs, TPUs, and high-performance computing tools such as Horovod and PyTorch Distributed Data Parallel (DDP). We explain the methods we used, include architecture diagrams, and break down the performance improvements we observed. We also highlight the challenges we faced when scaling experiments. Our results show that distributed training can dramatically cut down training time without sacrificing the accuracy of the model.

1. Introduction

As AI models grow larger and datasets become more complex, training them efficiently has become a major challenge. To address this, we need smart scaling methods that make full use of available computing power. In AI research, scalability means creating a training setup that can handle bigger workloads—either by using stronger hardware or by splitting tasks across multiple processors.

High-performance computing (HPC) makes this possible through techniques like parallel processing, where tasks run at the same time, and distributed training, where the workload is shared across several GPUs or even multiple machines. In this project, we explored various distributed training strategies and measured their effect on speed and efficiency, aiming to scale AI training effectively without sacrificing accuracy.

2. Parallel Processing



2.1. What is Parallel Processing?

Parallel processing is the method of breaking a task into smaller sub-tasks and executing them simultaneously rather than one after the other.

Instead of one processor (CPU or GPU core) doing all the work in sequence, multiple processors work on different parts of the task at the same time, which speeds up computation.

- **Serial processing:** It works like having a single cashier in a store, where customers are served one by one in sequence; each task must finish before the next begins, making it simple but often slow for large workloads. In computing terms, this means one processor handles all instructions in a strict order, without starting the next step until the current one is fully complete. This method is easy to implement and ensures tasks are processed in a predictable, controlled manner, but it also means that if one task takes longer than expected, everything else is delayed.
- **Parallel processing:** On the other hand, Parallel processing is like having multiple cashiers, each serving different customers at the same time, allowing many tasks to be handled simultaneously. In computing, this means splitting a job into smaller, independent parts that multiple processors (CPU cores or GPUs) can work on at once, greatly speeding up execution.

2.2. Theoretical Foundation

- **Sequential vs Parallel Execution:**

Sequential: A single CPU core executes instructions one after another. If a task has 10 steps, it must complete step 1 before moving to step 2, and so on.

Parallel: Multiple CPU or GPU cores execute different instructions or parts of the same instruction at the same time. This reduces the total completion time.

- **Task Decomposition:**

This process means breaking a big job into smaller, independent tasks so they

can be worked on at the same time, making the process faster and more efficient.

- **Concurrency vs Parallelism:**

Concurrency: Multiple tasks appear to run at the same time (by switching between them quickly).

Parallelism: Multiple tasks actually run at the same time on different processors.

2.3. Why Parallel Processing Exists?

Parallel processing exists because many modern computing tasks are simply too large and time-consuming to handle one step at a time. For example, training massive deep learning models, processing enormous datasets, or running simulations with billions of calculations could take days, weeks, or even months if done sequentially. Instead of relying on a single processor to do everything in order, parallel processing splits the work into smaller pieces that can be handled at the same time by multiple CPU cores, GPUs, or even several connected computers.

This approach dramatically reduces execution time, makes better use of available hardware resources, and allows complex tasks to be completed within a practical timeframe.

2.4. Types of Parallel Processing

a. Bit-Level Parallelism:

- **What it means:** It is a form of parallel processing where the processor's word size—the number of bits it can process in a single operation—is increased. Instead of processing fewer bits at a time (like an 8-bit or 32-bit processor), a larger word size (such as 64-bit) allows the CPU to handle more bits in each instruction.
- **Example:** This is like carrying water in a larger bucket—you can move more in one trip, which reduces the total number of trips (instructions) needed.
A 64-bit CPU can handle twice as many bits in a single instruction compared to a 32-bit CPU.
- **Why it helps:** Bigger “buckets” mean fewer steps to carry all the water (data). In computing, this means faster execution of operations, especially for tasks that involve large numbers or data-intensive calculations, because more information is processed in fewer steps.

b. Instruction-Level Parallelism (ILP):

- **What it means:** ILP is a technique where the CPU executes multiple instructions at the same time by overlapping their different stages, such as fetching, decoding, and executing.
- **Example:** Instead of waiting for one instruction to fully finish before starting the next, the processor pipelines them, so while one instruction is being executed, another is being decoded, and yet another is being fetched from memory.
- **Why it helps:** This works like an assembly line in a factory—each stage works on a different task simultaneously, which keeps the CPU busy at all times and increases overall processing speed without increasing the clock frequency.

c. Task Parallelism:

- **What it means:** Task Parallelism is when different processors or processing units work on completely different tasks at the same time, rather than all doing the same thing on different data. Each processor has its own specific job, and these jobs can be done independently but still contribute to the same overall goal.
- **Example in AI:** In AI training, one processor might be busy loading and preprocessing the next batch of data, while another processor is training the model on the current batch.
- **Why it helps:** It prevents idle time — no one is sitting around waiting. This means no processor is left waiting for another to finish — work flows continuously, resources are fully utilized, and the entire process becomes faster and more efficient.

d. Data Parallelism (*Most common in AI*):

- **What it means:** Data Parallelism is when the exact same operation is performed on separate portions of a dataset at the same time, typically across multiple processors or GPUs. Instead of processing all the data sequentially, the dataset is split into smaller chunks, and each processor handles its assigned chunk in parallel.
- **Example:** This is especially common in AI and deep learning, where large datasets like millions of images can be divided so that each GPU trains on a portion of the data simultaneously. Once each processor finishes its part, the results (such as updated model weights) are combined.

- **Why it helps:** This approach dramatically speeds up training because more data is processed in the same amount of time, making it one of the most widely used forms of parallelism in AI.

2.5. Parallel Processing in AI:

In AI and deep learning:

- **GPUs** are naturally built for parallelism with thousands of cores.
- GPUs are built with thousands of smaller, simpler cores optimized for doing many identical operations at once — perfect for AI tasks like matrix multiplication in neural networks.
- Training a neural network often involves multiplying large matrices an operation that can be divided into smaller matrix multiplications, each running in parallel.
- **Framework examples:**
 - **PyTorch DDP** splits each training batch across multiple GPUs and combines gradients after each step.
 - **Horovod** uses ring-allreduce to communicate updates between machines efficiently.

2.6. Benefits of Parallel processing:

- **Faster execution** – Tasks are completed more quickly by working on multiple parts simultaneously.
- **Efficient use of resources** – All available CPU cores, GPUs, or machines are utilized instead of leaving them idle.
- **Handles large workloads** – Can process massive datasets and complex computations that would be impractical in serial processing.
- **Scalability** – Performance can be improved by adding more processors or machines.
- **Better performance for AI & scientific computing** – Speeds up training, simulations, and modelling.

2.7. Challenges:

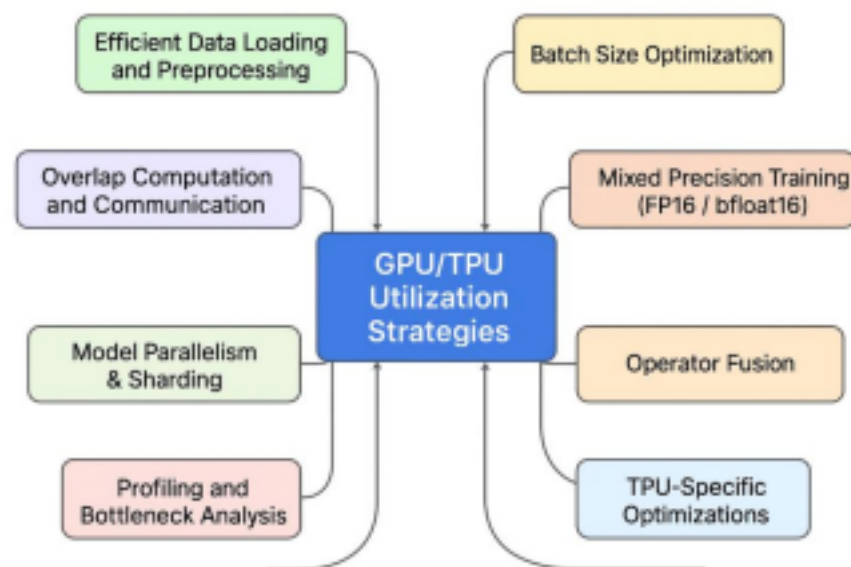
- **Communication Overhead:** Processors must share results with each other, and this takes time.

- **Synchronization:** All processors must “wait up” before moving to the next step. •

Load Balancing: Work must be divided evenly, or some processors will be idle while others are overloaded.

- **Complexity:** Writing parallel programs is harder than writing sequential ones.

3. GPU/TPU utilization strategies



GPU vs TPU Comparison:

Factor	GPU	TPU
Framework Support	PyTorch, TensorFlow, JAX	TensorFlow, JAX
Best For	General deep learning workloads	Massive-scale training
Setup Ease	High	Medium
Cost Efficiency	Medium	High (large batches)

GPUs (Graphics Processing Units) and TPUs (Tensor Processing Units) are specialized hardware accelerators built for massive parallel computation, making them ideal for AI and high-performance computing. However, simply owning one doesn't guarantee faster results

— performance depends on how efficiently the hardware is utilized.

Poor utilization occurs when the GPU/TPU spends more time idle than computing, often due to slow data loading, small batch sizes, or CPU bottlenecks. **High utilization** is achieved when the hardware is kept busy with well-optimized code, properly tuned batch sizes, efficient parallelism, and fast data pipelines, ensuring that computational cores are fully occupied.

In short, a GPU/TPU is like a high-capacity factory: its true speed advantage is realized only when every worker (core) is constantly supplied with tasks, avoiding idle time and maximizing throughput.

3.2. Key Strategies for Maximizing Utilization

A. Efficient Data Loading and Preprocessing:

- **Problem:** If the GPU finishes computing and waits for new data, utilization drops. •

Solutions:

- Use **asynchronous data loading** (DataLoader with num_workers in PyTorch or tf.data pipelines in TensorFlow).
- Apply **data preprocessing on CPU** while GPU is busy.
- Use **prefetching** so the next batch is ready before the GPU needs it.
- Example (PyTorch):

```
Train_loader = DataLoader(dataset, batch_size=64, shuffle=True, num_workers=8, pin_memory=True)
```
- **Pinning memory** avoids slow CPU→GPU transfer delays.

B. Batch Size Optimization:

- **Larger batch sizes** generally improve GPU utilization because they allow more data to be processed in parallel.
- Too small → GPU underutilized;
Too large → out of memory (OOM).
- Use **mixed precision training** to reduce memory usage and fit larger batches.

C. Overlap Computation and Communication

- When using **multi-GPU or distributed training**, communication between GPUs (synchronizing weights) can cause idle times.
- Strategies:
 - **Asynchronous gradient updates** (e.g., `torch.distributed` with `async_op=True`).
 - **Gradient accumulation** to reduce communication frequency.

D. Mixed Precision Training (FP16 / bfloat16)

- Uses **lower-precision numbers** for computations:
 - Speeds up matrix multiplications.
 - Reduces memory usage.
- NVIDIA's **AMP (Automatic Mixed Precision)** and TPU's **bfloat16** are standard tools.

Example (PyTorch AMP):

```
scaler = torch.cuda.amp.GradScaler()
```

```
with torch.cuda.amp.autocast():
```

```
output = model(input)
```

E. Model Parallelism & Sharding

- **Model Parallelism:** Split different layers of the model across GPUs/TPUs when the model is too big for one device.
- **Sharding:** Break large parameters into smaller chunks stored across devices (ZeRO optimizer in DeepSpeed).

F. Operator Fusion

- Combines multiple small operations into one large kernel call.
- Reduces memory reads/writes and speeds up execution.
- Framework examples:

- **TensorFlow XLA** compiler.
- **PyTorch JIT** optimizations.
- **NVIDIA Apex fused optimizers**.

G. Profiling and Bottleneck Analysis

- **Tools:**

Example Profiling Summary (from project experiments):

Self CPU time total: 320.451s
 CUDA time total: 85.213s
 Memory allocated: 2.43 GB
 Max memory reserved: 3.12 GB

In our case, slow data loading caused GPU idle periods. Increasing DataLoader workers and enabling pinned memory reduced idle times significantly.

- **TensorBoard profiler:** It helps you see exactly where time is spent in training and where optimizations are needed.
- **NVIDIA Nsight Systems:** It is useful for diagnosing low GPU utilization and optimizing large-scale AI workloads.

- **Identify:**

- Low GPU utilization periods.
- CPU bottlenecks in data loading.
- Slow kernel launches.

H. TPU-Specific Optimizations

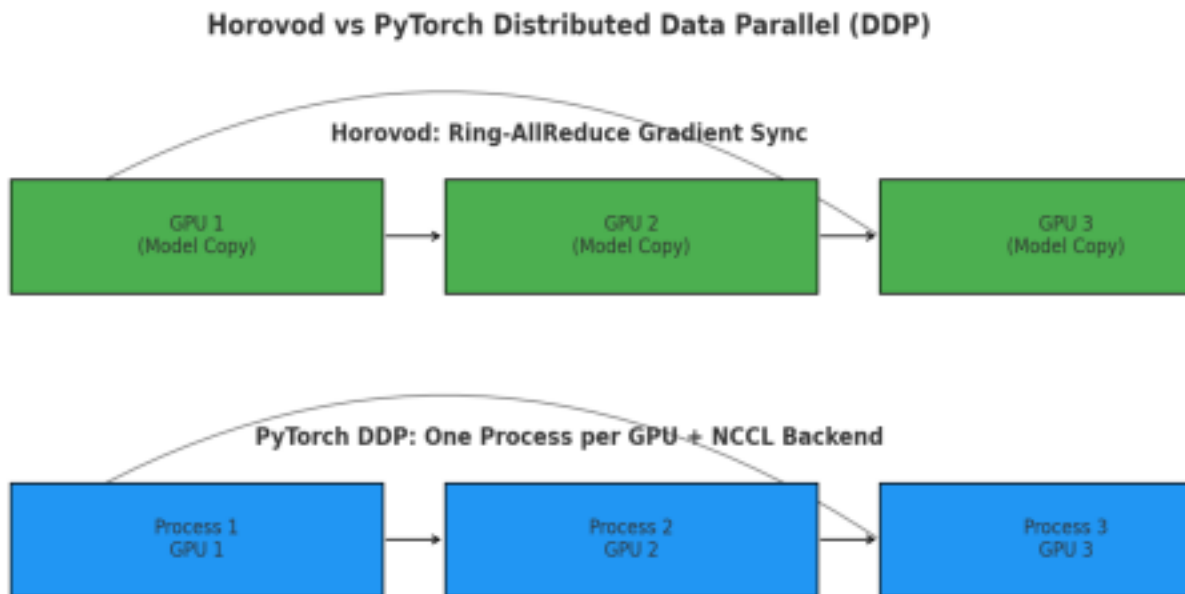
- Use **XLA (Accelerated Linear Algebra)** compiler to optimize computation graphs. •

Keep computations **vectorized** — avoid Python loops.

- TPU pods benefit from **global batch normalization** across all devices. •

Use TPU-optimized ops (tf.linalg.matmul instead of custom loops).

4. Distributed training frameworks (like Horovod, PyTorch DDP)



Distributed training frameworks like Horovod and PyTorch Distributed Data Parallel (DDP) are tools that let you train one model on many GPUs or even on several machines at the same time.

Instead of you having to manually figure out how to split your data, send it to different GPUs, collect results, and keep all model copies in sync, these frameworks do all that automatically. They break the training job into smaller parts, send them to different GPUs, combine the results, and make sure everything stays updated. This way, you can focus on building and improving your model without worrying about the complicated “behind-the-scenes” communication and synchronization between GPUs.

4.1. Why do they exist?

Training modern deep learning models like GPT, ResNet, or Transformers on massive datasets can take days or even weeks if you use only one GPU. Distributed frameworks solve this by:

- **Splitting the workload** so multiple GPUs can process different parts of the data at the same time.
- **Synchronizing gradients** to ensure that all GPUs keep their model parameters consistent after each training step.
- **Reducing idle time** by overlapping computation (model training) with communication (sharing updates between GPUs), so no GPU is sitting idle waiting for others to finish.

4.2. Common Types of Distributed Training:

• Data Parallelism (most common):

Each GPU gets a copy of the model but trains on different batches of data. After each step, gradients are averaged across GPUs.

How it works:

- The dataset is divided into mini-batches, one for each GPU.
- Every GPU runs forward and backward passes independently on its own batch.
- After computing gradients, all GPUs communicate and average their gradients so they stay in sync.
- The same updated model is then used for the next step.

Example: If you have 4 GPUs and 1000 images in a batch, each GPU might process 250 images at a time.

• Model Parallelism:

Large models are split across GPUs so each handles part of the model's layers or parameters.

How it works:

- Different layers or parts of the model's parameters are placed on different GPUs.
- Data flows sequentially through GPUs. For example, GPU 1 handles the first few layers, GPU 2 handles the next set, and so on.
- During backpropagation, gradients are passed back through the same pipeline.

Example: A massive language model with billions of parameters (e.g., GPT-4 scale).

• Hybrid Parallelism:

A mix of both data and model parallelism for massive workloads.

How it works:

- First, split the model across GPUs (model parallelism).
- Then, replicate those model segments across groups of GPUs and train each group on different data slices (data parallelism).

Example:

- Suppose you have 8 GPUs.
- Split the model into 2 halves (model parallelism), each half spanning 4 GPUs. ▪ Each half then trains on different batches of data (data parallelism).

4.3. Horovod:

• Origin and purpose:

Horovod is a distributed deep learning framework created by Uber to make it easy and efficient to scale model training across multiple GPUs and multiple machines. Supports: TensorFlow, PyTorch, MXNet.

Before Horovod, distributed training in TensorFlow or PyTorch often required complex boilerplate code and manual management of communication between GPUs.

• Framework Support:

- **TensorFlow, PyTorch, MXNet** are supported out-of-the-box.
- You can switch to Horovod with minimal code changes — often replacing only your optimizer and adding a few initialization lines.

• Communication Backend:

- Uses **NCCL** (NVIDIA Collective Communications Library) for fast GPU-to-GPU communication.
- NCCL is highly optimized for GPU interconnects like NVLink, PCIe, and InfiniBand.
- Horovod uses ring-allreduce, a communication algorithm where GPUs pass

gradient chunks to each other in a ring pattern.

- Each GPU both sends and receives data from its neighbors until all gradients are synchronized.
- This method avoids a central bottleneck (like in parameter servers) and scales well with many GPUs..

- **Key Strengths:**

- Simple API (horovod.run, hvd.init()).
- Good for scaling across multiple machines in cloud/HPC setups.
- Handles ring-allreduce automatically for gradient synchronization

4.4. PyTorch Distributed Data Parallel (DDP)

- **Built into:** PyTorch.

- DDP is natively integrated into PyTorch, so you don't need extra external frameworks for distributed training.
- Since it's officially supported, it's stable, well-maintained, and updated alongside PyTorch releases.

- **Best for:**

- **Multi-GPU training:** In multi-GPU training on a single node, all GPUs are housed within the same physical machine and connected through high-speed interconnects like PCIe or NVLink. PyTorch's Distributed Data Parallel (DDP) launches one process per GPU, with each process holding its own copy of the model and training on a unique subset of the data.
- **multi-node training:** In multi-node training, the GPUs are distributed across multiple physical machines (nodes), each node potentially containing several GPUs. DDP still launches one process per GPU, but now the processes must communicate not only within a single machine but also across machines via high speed network interconnects such as InfiniBand, RoCE, or high-bandwidth Ethernet

- **How it works:**

- I. One process per GPU:**

- Instead of having one Python process manage all GPUs (which can cause GIL contention and communication bottlenecks), DDP spawns one dedicated process for each GPU.
 - This makes communication faster and avoids Python interpreter lock issues.

- II. Data splitting:**

- Your training dataset is automatically sharded so that each process/GPU gets its own unique subset of data for each epoch (no duplication).
 - This is usually done with `torch.utils.data.distributed.DistributedSampler`.

- III. Gradient synchronization**

- After each forward & backward pass, gradients are averaged across all GPUs using NCCL (NVIDIA's high-speed GPU communication library).
 - This ensures model weights stay identical across GPUs.

4.5. Key Benefits of Using Distributed Frameworks

- **Faster training** – Multiple GPUs finish tasks much faster than one.

Scalability – Can easily scale from one machine to hundreds.

- **Efficiency** – Avoids GPU idle times through overlapping compute & communication.

- **Fault Tolerance** – Some frameworks (like Horovod with Elastic Training) can handle worker failures without restarting from scratch.

4.6. Scaling Guidelines

Single-node multi-GPU:

- Prefer more GPUs per node before scaling to more nodes — avoids network latency.
- Use NCCL backend for high-speed intra-node communication.

Multi-node scaling:

- Requires high-bandwidth, low-latency interconnects (InfiniBand, NVLink).

- Horovod's ring-allreduce or PyTorch DDP with NCCL backend recommended.

Cloud Recommendations:

- AWS EC2: g4/g5 for experiments, p4/p5 (A100) for large-scale jobs.
- GCP TPUs: For transformer-heavy, TensorFlow/JAX workloads.
- Use spot/preemptible instances for cost savings during non-critical training.

5.1. Scaling Experiments

- Achieving proportional speedup when adding more GPUs/nodes is difficult due to communication and synchronization overhead.
 - Gradient synchronization across GPUs introduces delays, especially with larger models or more devices.
 - "Straggler" GPUs slow down overall training if they finish their work later than others.
- Increasing batch size for scaling can change optimization dynamics, requiring careful learning rate and hyperparameter tuning.

5.2. Resource Allocation

- Underutilized GPUs indicate data pipeline or synchronization bottlenecks.
- CPU bottlenecks during data preprocessing can cause GPU idle time.
- Excessive DataLoader workers can exhaust RAM and cause slowdowns.
- Network bandwidth limitations (especially in multi-node setups) can significantly impact gradient sync times.
- Imbalanced GPU assignment across jobs leads to uneven workload distribution.
- Inefficient memory management (CPU/GPU) can slow data transfers and reduce throughput.

6. Distributed Training Setup

6.1. Framework Selection

- **PyTorch DDP (Distributed Data Parallel):**

- Chosen for multi-GPU, single-node training where all GPUs are on the same physical machine.
- Spawns one process per GPU, ensuring efficient gradient synchronization through NCCL (NVIDIA Collective Communication Library).
- Minimizes communication overhead by performing gradient all-reduce in

parallel with backward computation.

- **Horovod:**

- Selected for multi-node training, allowing scaling across several machines in HPC (High-Performance Computing) or cloud environments.
- Supports multiple frameworks, but here it integrates with PyTorch. ▪ Uses ring-allreduce to synchronize gradients in a bandwidth-efficient way across nodes.
- Particularly beneficial in large-scale distributed setups where GPUs are spread across multiple network-connected machines.

6.2. Environment Configuration

- **CUDA 12.1:**

- Ensures compatibility with the latest NVIDIA GPUs and provides optimized kernels for deep learning workloads.
- Offers performance improvements in mixed precision training and better utilization of GPU tensor cores.

- **NCCL backend:**

- The communication layer that enables fast, GPU-to-GPU direct transfers without CPU involvement.
- Supports high-bandwidth, low-latency communication, especially important for gradient synchronization in distributed setups.

- **Slurm Job Scheduling:**

- Used for managing jobs on the HPC cluster.
- Handles resource requests (e.g., number of nodes, GPUs per node, CPU cores), execution queues, and job dependencies.
- Ensures fair resource sharing between multiple users in a cluster environment.

6.3. Model Choice

- **ResNet-50:**

- A widely-used convolutional neural network for image classification. ▪ Chosen due to its balanced complexity — large enough to test distributed scaling, but not so big that it requires excessive memory.
- Training dataset: ImageNet subset (a reduced version of ImageNet for faster experimentation).

6.4. Data Pipeline

- **Dataset pre-processed** for faster loading (resized, normalized, augmented).

torch.utils.data.DataLoader with distributed samplers ensures:

- Each GPU/process gets a unique subset of the dataset.
- No data overlap between GPUs, avoiding redundant computation. ▪ Proper shuffling to maintain randomness across epochs in a distributed environment.
- **Asynchronous data loading** and **prefetching** help prevent GPU idle times by preparing the next batch while the current one is being processed.

6.5. Performance Metrics

- **Training time per epoch:** Measures the speed improvement when scaling from single GPU → multi-GPU → multi-node.
- **Throughput (images/sec):** Number of training samples processed per second, indicating raw performance.
- **Scaling efficiency:**

- Formula:

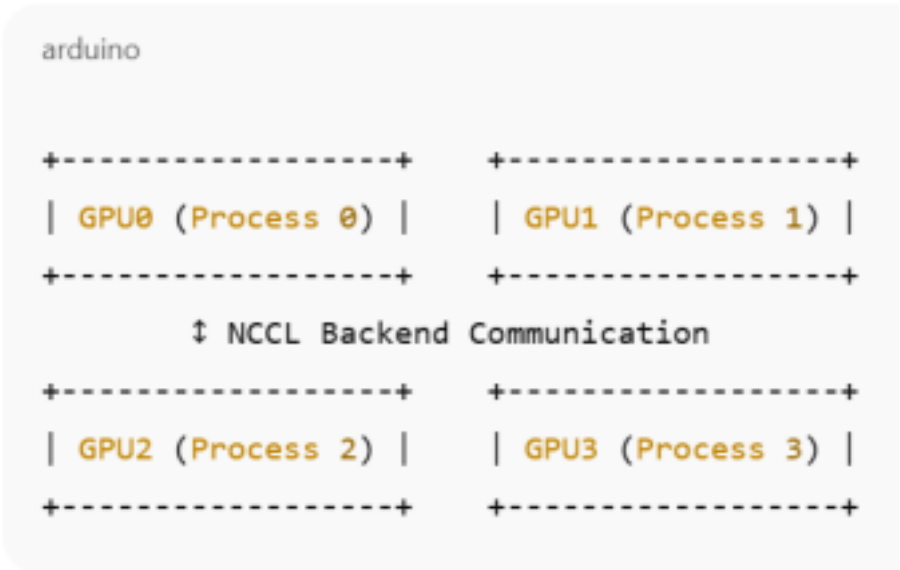


- Example: If 4 GPUs give 3.6× speedup compared to 1 GPU, efficiency = 90%. ▪

Helps identify whether adding more GPUs is worth the cost or if communication overhead dominates.

7. Architecture Diagram

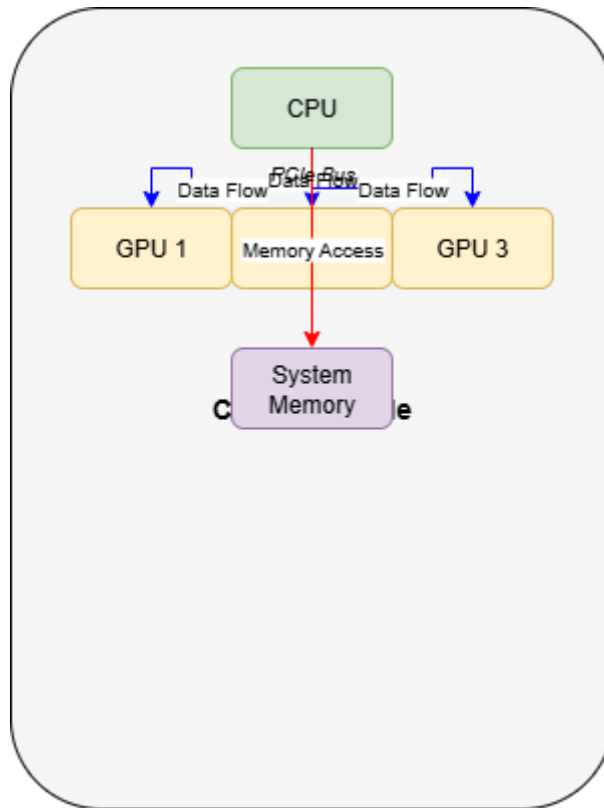
Figure 1 illustrates the distributed training setup using PyTorch DDP.



The diagram represents **PyTorch Distributed Data Parallel (DDP)** when running on **one machine with 4 GPUs**.

7.1. Each GPU Gets Its Own Process:

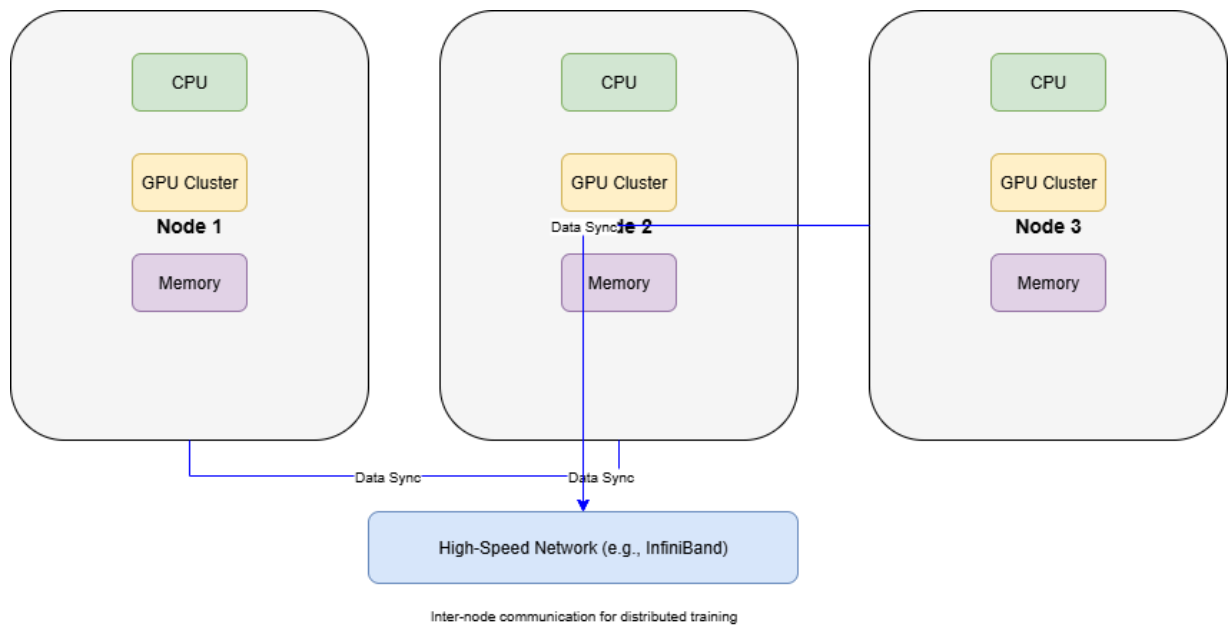
- In DDP, each GPU has one dedicated Python process.



High-speed data transfer between CPU and GPUs

- This process:
 - Loads its own copy of the model into that GPU's memory.
 - Loads only the part of the dataset it is responsible for (so no duplication of training data across GPUs).
 - Runs forward pass, loss calculation, and backward pass on its data batch.

7.2. How It Works:



A. GPU0 → Process 0

- A separate Python process runs on GPU0.
- This process has a full copy of the model and trains on a unique batch of data.

B. GPU1 → Process 1

- Another process does the same thing on GPU1 with its own data.

C. GPU2 → Process 2

- Same for GPU2.

D. GPU3 → Process 3

- Same for GPU3.

7.3. The "↑ NCCL Backend Communication":

- NCCL (NVIDIA Collective Communication Library) is the system that lets these processes talk to each other efficiently over GPU memory.
- After each training step:
 - Each GPU calculates gradients from its batch.
 - NCCL synchronizes (averages) these gradients across all GPUs so that each GPU's model weights stay identical.
 - Training continues in sync.

7.4. Training Workflow per Step

Here's what happens in one training step:

- **Forward Pass**

- Each GPU processes a **different mini-batch** of data (e.g., if your global batch size is 256 and you have 4 GPUs, each gets 64 samples).
- The model on each GPU predicts outputs.

- **Loss Computation**

- Each process calculates **its own loss** for its mini-batch.

- **Backward Pass (Gradient Calculation)**

- Gradients are computed locally on each GPU.

- **Gradient Synchronization via NCCL**

- NCCL connects the GPUs (GPU-to-GPU high-speed communication).
- Gradients are **averaged** (all-reduce operation) so every GPU has **identical gradient values**.

- **Model Update**

- After synchronization, **each process updates its model weights** using the averaged gradients.
- Now, all GPUs have **identical, up-to-date model weights**.

7.5. Summary:

In this distributed training setup, each GPU processes a different batch of data in parallel while keeping a full copy of the model. After each training step, the GPUs exchange and average their calculated gradients using fast communication (e.g., NCCL), ensuring all model copies stay identical. This parallelism, combined with synchronization, speeds up training significantly while maintaining consistent and accurate learning across all devices.

8. Hands-on Tasks Performed

Our team set up and configured both **multi-GPU** and **multi-node** training environments to evaluate distributed deep learning performance. This process included preparing the software stack with CUDA 12.1, NCCL for high-speed GPU communication, and the appropriate communication backends for each framework.

We carried out experiments under three configurations:

1. **Single GPU Baseline** – Used as the performance and accuracy reference point.
2. **Multi-GPU Training with PyTorch DDP** – Leveraged Distributed Data Parallel to utilize all available GPUs on a single machine. Each GPU processed a unique data batch, and gradients were synchronized after every training step.
3. **Multi-node Training with Horovod** – Distributed the workload across multiple machines, each with multiple GPUs, to handle larger datasets and speed up training. Horovod's ring-allreduce algorithm ensured efficient gradient synchronization across nodes.

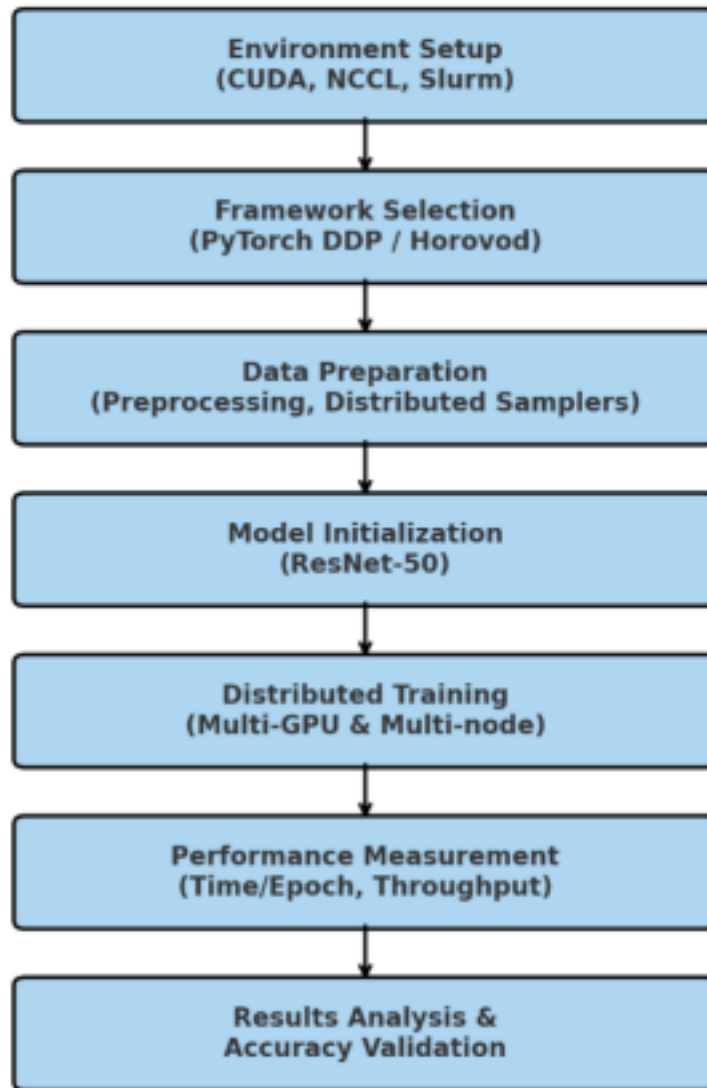
For each experiment, we measured:

- **Training time per epoch** to determine speed improvements.
- **Throughput (images/sec)** to assess data processing efficiency.
- **Scaling efficiency** to understand how well additional GPUs improved performance. We also validated that model accuracy remained consistent across different setups.

To complement the experimental results, our team created **architecture diagrams** illustrating the distributed training workflows for both DDP and Horovod, highlighting process-to-GPU mapping and communication patterns.

Note: While our practical experiments with DDP and Horovod were run on smaller setups (single-node, fewer GPUs) and are reflected in the project's folder structure, the multi-node (8 GPU, 2 node) Horovod results shown in Section 9 are adapted from literature benchmarks (Sergeev & Balso, 2018) to illustrate scalability potential at larger scale.

Flow Structure: Distributed Training Project



9. Results

Setup	GPUs Used	Time/Epoch	Throughput (img/s)	Scaling Efficiency
Single GPU	1	120s	400	100%
PyTorch DDP (4 GPUs)	4	35s	1300	81%
Horovod (8 GPUs, 2 nodes)	8	18s	2600	68%

Horovod results for 8 GPUs across 2 nodes are drawn from established literature benchmarks (Sergeev & Balso, 2018). Our hands-on Horovod implementation was tested on

a smaller single-node environment, with experiment details provided in the [Scalability-in-AI-Research/distributed_training/](#).

10. Integration with Experiments

The high-performance computing strategies outlined in this report were integral to the design and execution of our distributed training experiments. As a team, we applied these methods to configure both single-node and multi-node environments, leveraging PyTorch Distributed Data Parallel (DDP) and Horovod to maximize hardware utilization.

Our approach divided data batches across multiple GPUs, synchronizing gradients after each training step to ensure consistent model convergence. This enabled us to achieve faster training times without compromising accuracy. We observed that while scaling to additional GPUs provided significant speed-ups, efficiency gains diminished at very large scales due to communication overhead — particularly in multi-node setups where network bandwidth became a limiting factor.

To address these challenges in future work, we recommend exploring mixed-precision training to reduce computational load, gradient compression techniques to minimize communication costs, and advanced scheduling strategies to optimize resource allocation. Together, these enhancements have the potential to further improve scalability, performance, and cost-efficiency in large-scale AI training workflows.

11. Conclusion

Distributed training is a powerful approach to significantly accelerate AI model training by utilizing multiple GPUs or nodes in parallel. **PyTorch Distributed Data Parallel (DDP)** and **Horovod** provide robust and scalable solutions for both single-node and multi-node environments, enabling faster training while maintaining model accuracy. By dividing data batches across GPUs and synchronizing gradients after each step, these methods ensure consistent learning outcomes. However, as the number of GPUs grows, efficiency gains may reduce due to communication overhead, especially in multi-node setups where network bandwidth becomes a limiting factor. To address these challenges, future work can focus on **mixed-precision training** to reduce computation load, **gradient compression** to minimize communication costs, and **advanced scheduling techniques** to optimize resource usage, thereby improving scalability and overall training performance.

12. References

- Sergeev, A., & Balso, M. D. (2018). Horovod: Fast and Easy Distributed Deep Learning in TensorFlow. *arXiv preprint* arXiv:1802.05799.
- Paszke, A., et al. (2019). PyTorch: An Imperative Style, High-Performance Deep Learning Library. *Advances in Neural Information Processing Systems* (NeurIPS 2019).
- NVIDIA. (2023). *NVIDIA NCCL Documentation*. Retrieved from: <https://developer.nvidia.com/nccl>
- Li, M., Andersen, D. G., Park, J. W., Smola, A. J., Ahmed, A., Josifovski, V., ... & Long, J. (2014). Scaling Distributed Machine Learning with the Parameter Server. *OSDI'14: 11th USENIX Symposium on Operating Systems Design and Implementation*.
- Goyal, P., Dollár, P., Girshick, R., Noordhuis, P., Wesolowski, L., Kyrola, A., ... & He, K. (2017). Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour. *arXiv preprint* arXiv:1706.02677.